

Figure 4. H3: safety-axis predictors vs. TruthfulQA, per (dataset, model).

Rows: harm_content_density sub-dimension (top), composite S (bottom). Columns: MC1 (left of each block), MC2 (right). Asterisk indicates 95% CI excludes zero. H3 predicts NEGATIVE correlations.

